

AI Product Readiness Checklist

Run an AI effort through this before it ships. Every item is a human decision, not a model setting.

- **Intent.** Can you write the user's goal in one sentence, in their words, with no mention of AI? (*Problem statement + job story.*)
- **Tool fit.** Can you name the tool you chose, the simpler option you ruled out, and which agent triggers the work meets? (*Tool decision record + agent-fit assessment.*)
- **Expectations.** Can a first-time user see what the system can do, how often it's wrong, and how sure it is about this answer? (*Capability-and-limits statement + confidence-display spec.*)
- **Oversight.** Can you list every system the AI can change, name the one accountable owner, and show the way to stop it? (*Access matrix + accountability record + stop-and-rollback runbook.*)
- **Context.** Have you listed the context the model needs, confirmed each item is retrievable at run time, and closed the gaps? (*Context inventory + gap list.*)
- **Language.** Does the prompt carry all four CARE parts, with at least one worked example? (*CARE template + worked-example set.*)
- **Evals.** Do you have a scored eval set of real cases, a numeric pass bar, and a rule that anything below it doesn't ship? (*Held-out eval set + rubric + release-gate policy.*)
- **Outcome.** Have you named the single metric, its baseline, the target, and the date you'll check it? (*Outcome statement + measurement plan.*)

If any box is unchecked, the gap is human and so is the fix. That's the good news.

From [Patrick Neeman](#) at [AI for UX Substack](#)